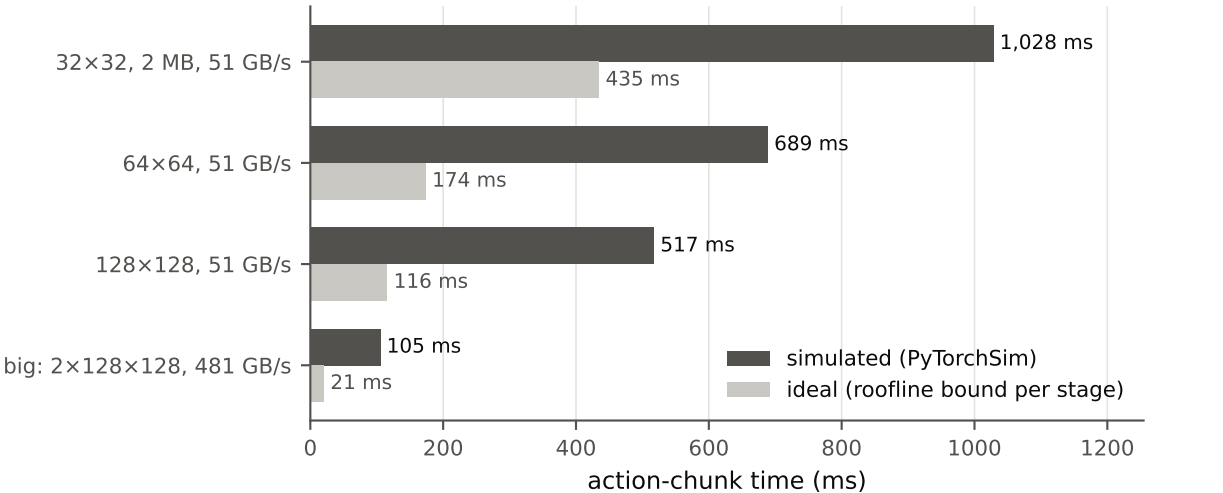
Chunk time vs array size (51.2 GB/s LPDDR5; big NPU for reference)
32×32, 2 MB, 51 GB/s
64×64, 51 GB/s
128×128, 51 GB/s
big: 2×128×128, 481 GB/s
1,028 ms
435 ms
689 ms
174 ms
517 ms
116 ms
105 ms
21 ms
simulated (PyTorchSim)
ideal (roofline bound per stage)
0
200
400
600
800
1000
1200
action-chunk time (ms)