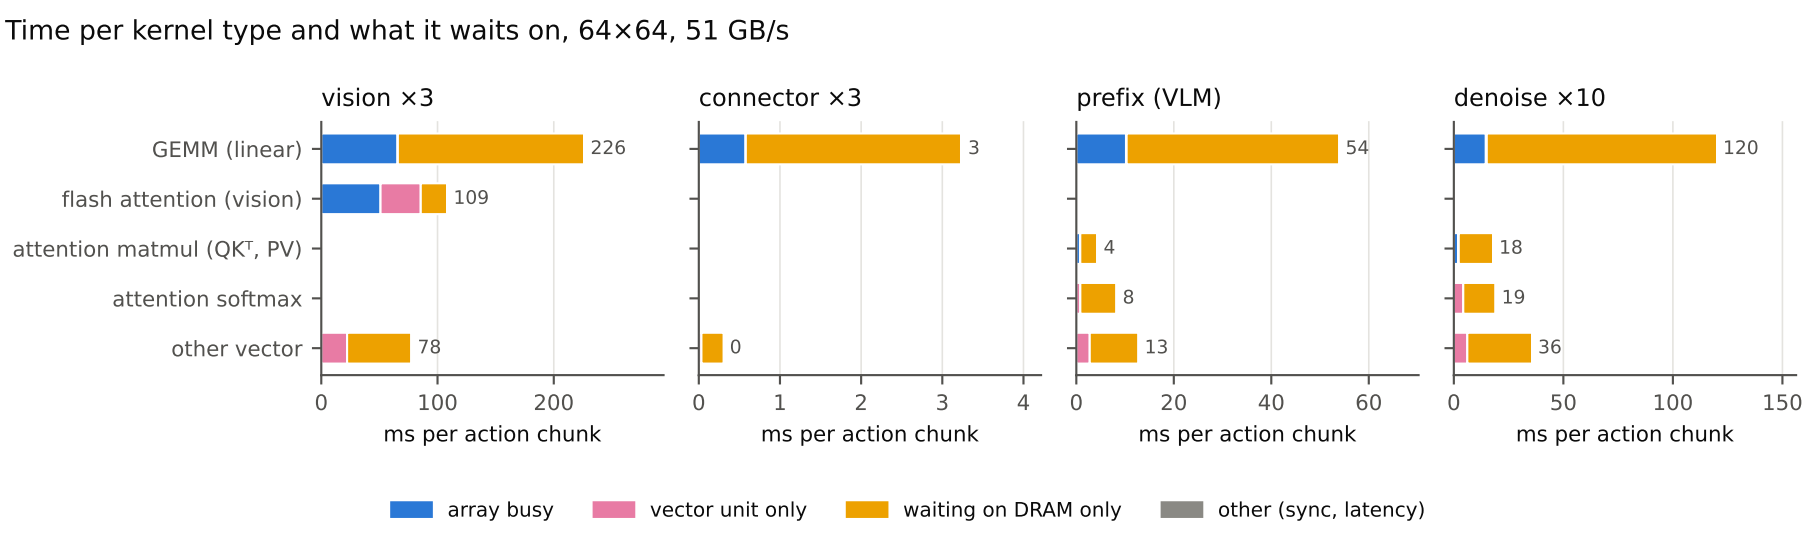
Time per kernel type and what it waits on, 64×64, 51 GB/s
vision ×3
connector ×3
prefix (VLM)
denoise ×10
GEMM (linear)
flash attention (vision)
attention matmul (QKᵀ, PV)
attention softmax
other vector
226
109
78
3
0
54
4
8
13
120
18
19
36
0
100
200
0
1
2
3
4
0
20
40
60
0
50
100
150
ms per action chunk
array busy
vector unit only
waiting on DRAM only
other (sync, latency)