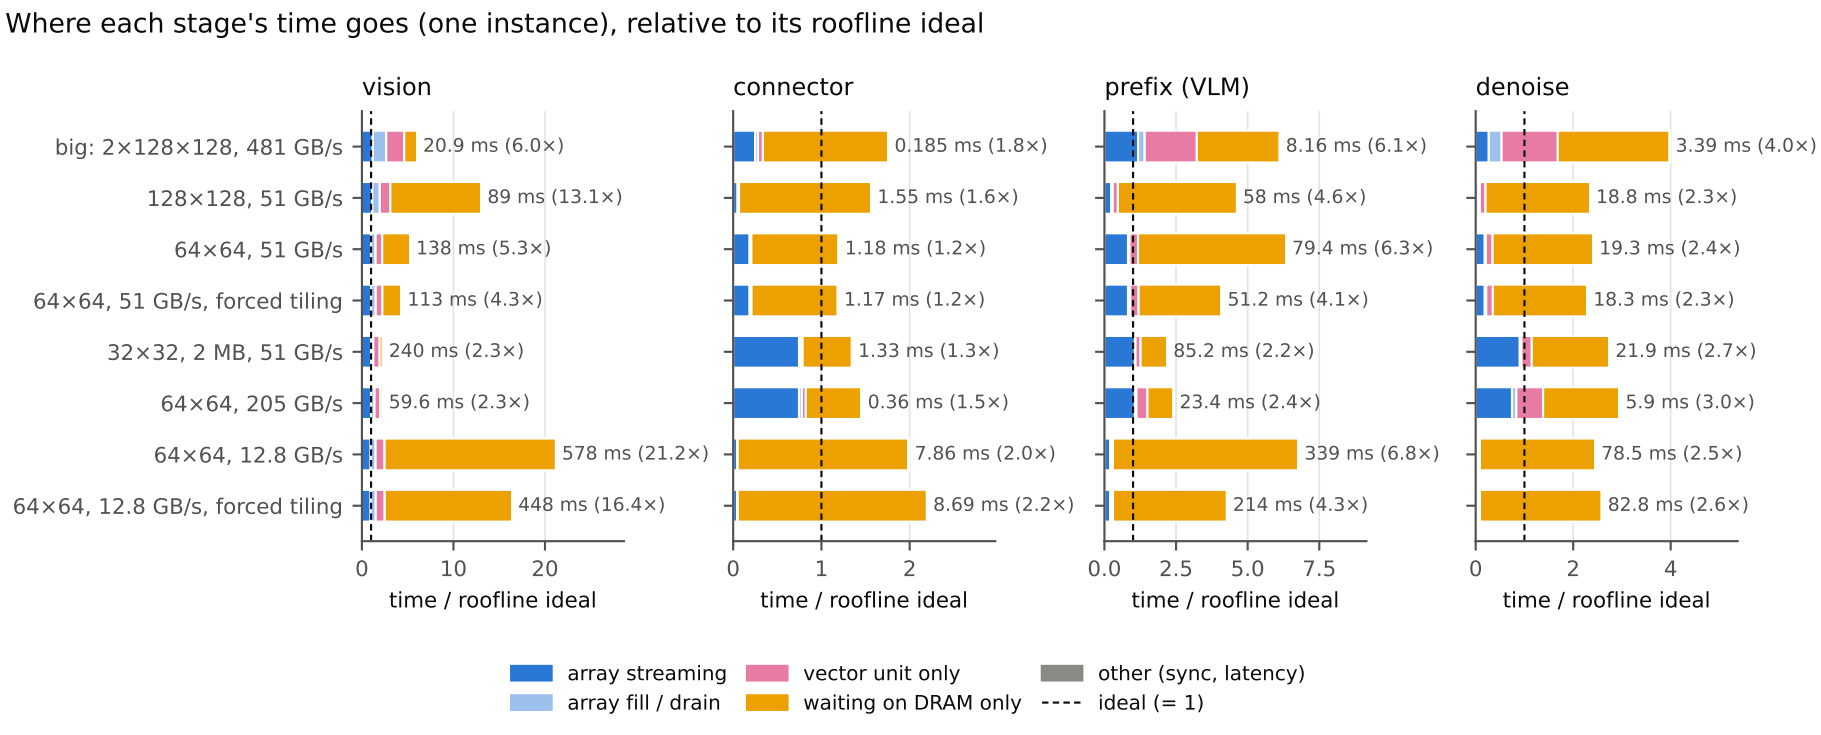
Where each stage's time goes (one instance), relative to its roofline ideal
vision
connector
prefix (VLM)
denoise
big: 2×128×128, 481 GB/s
128×128, 51 GB/s
64×64, 51 GB/s
64×64, 51 GB/s, forced tiling
32×32, 2 MB, 51 GB/s
64×64, 205 GB/s
64×64, 12.8 GB/s
64×64, 12.8 GB/s, forced tiling
20.9 ms (6.0×)
89 ms (13.1×)
138 ms (5.3×)
113 ms (4.3×)
240 ms (2.3×)
59.6 ms (2.3×)
578 ms (21.2×)
448 ms (16.4×)
0.185 ms (1.8×)
1.55 ms (1.6×)
1.18 ms (1.2×)
1.17 ms (1.2×)
1.33 ms (1.3×)
0.36 ms (1.5×)
7.86 ms (2.0×)
8.69 ms (2.2×)
8.16 ms (6.1×)
58 ms (4.6×)
79.4 ms (6.3×)
51.2 ms (4.1×)
85.2 ms (2.2×)
23.4 ms (2.4×)
339 ms (6.8×)
214 ms (4.3×)
3.39 ms (4.0×)
18.8 ms (2.3×)
19.3 ms (2.4×)
18.3 ms (2.3×)
21.9 ms (2.7×)
5.9 ms (3.0×)
78.5 ms (2.5×)
82.8 ms (2.6×)
0
10
20
0
1
2
0.0
2.5
5.0
7.5
0
2
4
time / roofline ideal
array streaming
array fill / drain
vector unit only
waiting on DRAM only
other (sync, latency)
ideal (= 1)