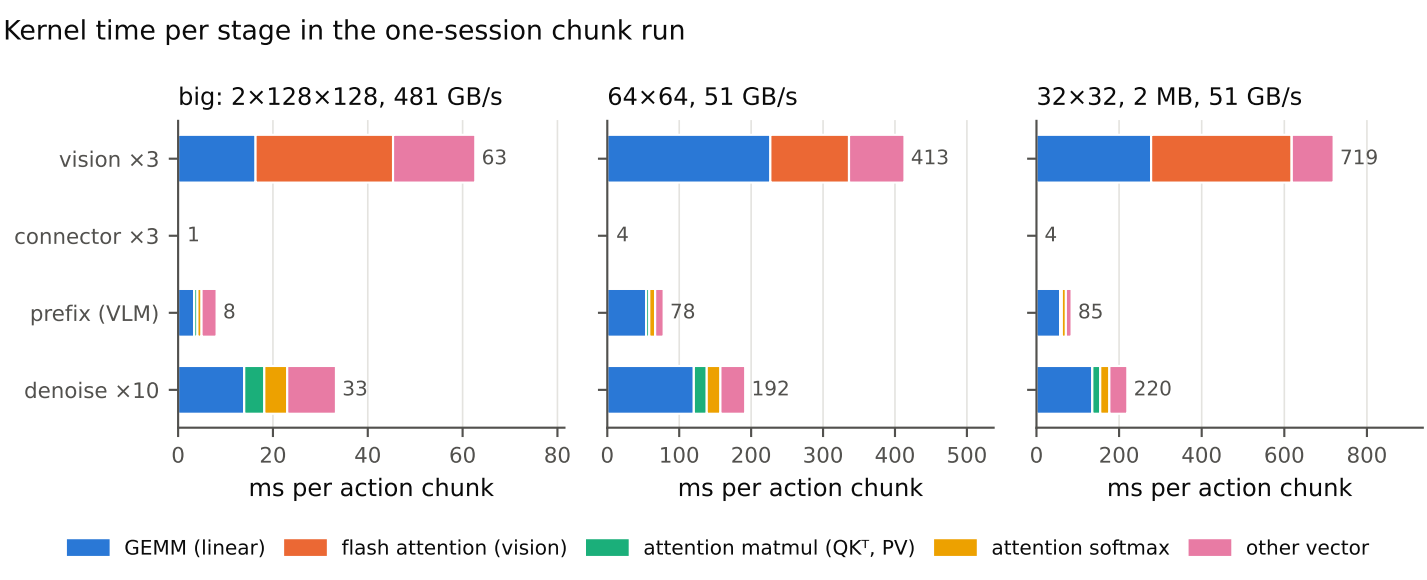
Kernel time per stage in the one-session chunk run
big: 2×128×128, 481 GB/s
64×64, 51 GB/s
32×32, 2 MB, 51 GB/s
vision ×3
connector ×3
prefix (VLM)
denoise ×10
63
1
8
33
413
4
78
192
719
4
85
220
0
20
40
60
80
0
100
200
300
400
500
0
200
400
600
800
ms per action chunk
ms per action chunk
ms per action chunk
GEMM (linear)
flash attention (vision)
attention matmul (QKᵀ, PV)
attention softmax
other vector