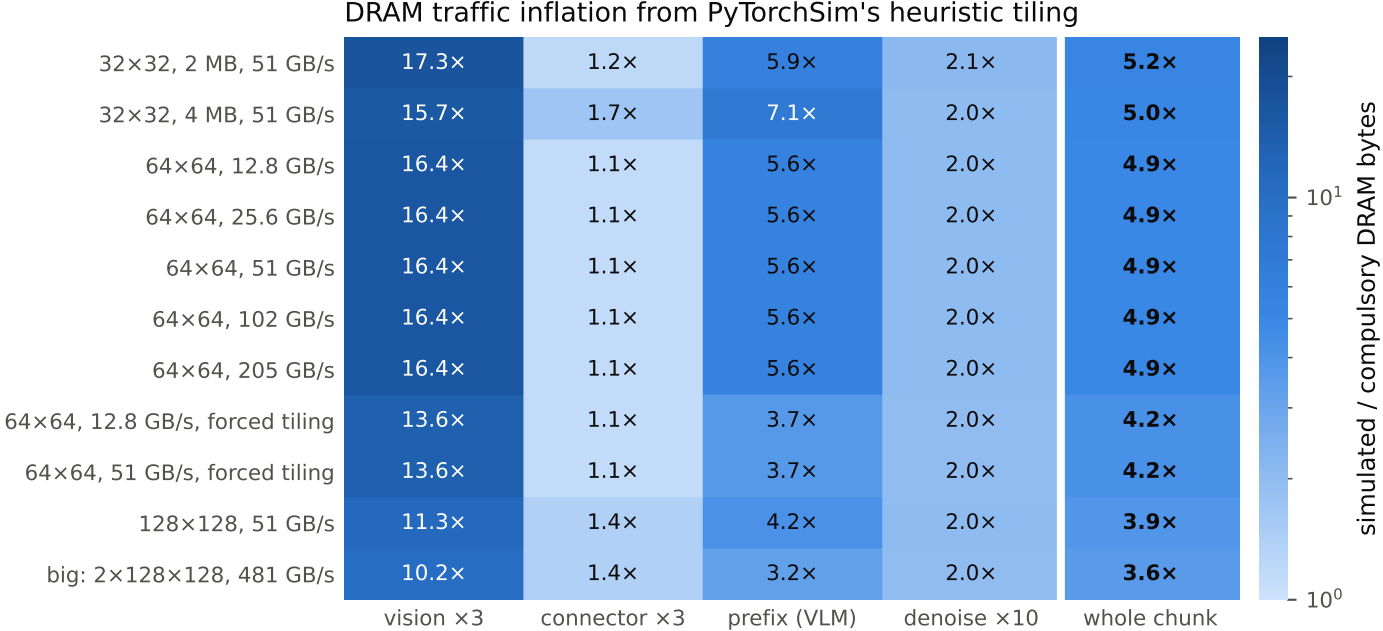
DRAM traffic inflation from PyTorchSim's heuristic tiling
32×32, 2 MB, 51 GB/s
32×32, 4 MB, 51 GB/s
64×64, 12.8 GB/s
64×64, 25.6 GB/s
64×64, 51 GB/s
64×64, 102 GB/s
64×64, 205 GB/s
64×64, 12.8 GB/s, forced tiling
64×64, 51 GB/s, forced tiling
128×128, 51 GB/s
big: 2×128×128, 481 GB/s
17.3× 1.2× 5.9× 2.1× 5.2×
15.7× 1.7× 7.1× 2.0× 5.0×
16.4× 1.1× 5.6× 2.0× 4.9×
16.4× 1.1× 5.6× 2.0× 4.9×
16.4× 1.1× 5.6× 2.0× 4.9×
16.4× 1.1× 5.6× 2.0× 4.9×
16.4× 1.1× 5.6× 2.0× 4.9×
13.6× 1.1× 3.7× 2.0× 4.2×
13.6× 1.1× 3.7× 2.0× 4.2×
11.3× 1.4× 4.2× 2.0× 3.9×
10.2× 1.4× 3.2× 2.0× 3.6×
vision ×3
connector ×3
prefix (VLM)
denoise ×10
whole chunk
simulated / compulsory DRAM bytes
10^1
10^0